

# Stan for the people

Two day introductory workshop  
on Bayesian modeling

McGill University  
January 25th 2019




[mc-stan.org](http://mc-stan.org)

## Part III

### Stan

- ▶ Stan is an expressive language for joint distributions.
- ▶ It automatically computes derivatives.
- ▶ It automatically performs inference algorithms.

# How Stan works



# How Stan works

- ▶ The Stan file specifies the joint distribution
$$p(\theta, z) = p(z|\theta)p(\theta)$$
- ▶ The input includes:
  - ▶ the data,  $z$
  - ▶ tuning parameters for the algorithm
- ▶ The output can include:
  - ▶ an approximate sample from the posterior distribution
  - ▶ summaries of the run which can help us diagnose problems.

# Inference algorithms in Stan

- ▶ Hamiltonian Monte Carlo (HMC)
- ▶ No-U Turn Sampler (NUTS)
- ▶ Automatic differentiation variational inference (ADVI)
- ▶ Lasso
- ▶ ...

We can manage the Stan file, the input, and the output using a scripting language, such as:

- ▶ R
- ▶ Python
- ▶ Julia
- ▶ The command line
- ▶ . . .

# Example 1: linear regression

The data generating process is:

$$Y \sim \text{Normal}(X\beta, \sigma^2)$$

Our goal is to estimate  $\theta = (\beta, \sigma^2)$ , based on the observation  $Z = (X, Y)$  and prior knowledge we have of  $\beta$  and  $\sigma$ .

► `data/linear.data.r`



# Example1: linear regression

As a prior, we use:

- ▶  $\beta \sim \text{Normal}(2.0, 1.0)$
- ▶  $\sigma^2 \sim \text{Gamma}(1.0, 1.0)$

which encode information from previously observed data.

# Writing the Stan file

We need a statement that specifies the log joint distribution.  
Recall:

$$p(\theta, z) = p(z|\theta)p(\theta)$$

Then:

$$\log p(\theta, z) = \log p(z|\theta) + \log p(\theta)$$

# Writing the Stan file

Stan retains certain C++ features:

- ▶ Variables need to be declared.
- ▶ Each statement must end with a semi-colon.

For example:

```
real x;
```

# Writing the Stan file

A Stan program is divided into coding blocks:

- ▶ data
- ▶ parameter
- ▶ model

# Writing the Stan file

```
data {  
  Declare the data that will be given as an input.  
}  
  
parameters {  
  Declare the parameters we want to sample.  
}  
  
model {  
  Compute the log joint distribution.  
}
```

# Writing the Stan file

```
model {  
  target += normal_lpdf(y | beta * x, sigma);  
  
  // or equivalently  
  
  y ~ normal(beta * x, sigma);  
}
```

# Writing the Stan file

Live demo.

# Diagnose

Are all (4) Markov chains exploring the target posterior?

Look at:


- ▶ the trace plots and the density plots
- ▶ the statistic  $\hat{R}$ .

The pairs plot can help us diagnose more difficult problems.



# Posterior predictive checks

- ▶ Recall Box's loop.
- ▶ Does our model accurately describe the data?



# Posterior predictive checks (ppc)

Given our posterior distribution for  $\theta$ , what kind of data,  $y_{\text{pred}}$ , do we generate?

Proposition:

*Each time we draw a sample,  $\theta^{(i)} = (\beta^{(i)}, \sigma^{(i)})$ , we will also simulate data, according to:*

$$y_{\text{pred}} \sim \text{Normal} \left( \mathbf{x} \beta^{(i)}, \sigma^{(i)} \right)$$

# Posterior predictive checks

To do this, we will use the `generated quantities` block.

Live demo.

# Improving the model

- ▶ The ppc suggest our model can improve with an intercept parameter.
- ▶ *Exercise: repeat the above procedure, but this time add a parameter  $\beta_0$ .*

# Remarks on the prior distribution

The prior distribution can encode:

- ▶ an existing posterior distribution
- ▶ theoretical information
- ▶ a regularization device (see Lasso, Ridge, etc.)
- ▶ any quantitative assumption.

The prior should be thought of as *part of the model*.

# Remarks on the prior distribution

Diagnostic tools:

- ▶ Just as we did posterior predictive checks, we can do *prior predictive checks*.
- ▶ Do our priors allow for all “reasonable” data configurations?
- ▶ Do they allow for “unreasonable” data configurations?

# Remarks on the prior distribution

Some helpful resources:

- ▶ <https://github.com/stan-dev/stan/wiki/Prior-Choice-Recommendations>
- ▶ Visualization in Bayesian workflow [Gabry et al., 2017]
- ▶ Towards a principled Bayesian workflow [Betancourt, 2018]



# General resources to use Stan

- ▶ The Stan user manual
- ▶ The (draft) Stan book ([https://mc-stan.org/docs/2\\_18/stan-users-guide/index.html](https://mc-stan.org/docs/2_18/stan-users-guide/index.html))
- ▶ The Stan forum (<http://discourse.mc-stan.org/>)

## Example 2: logistic regression


$$Y \in \{0, 1\}$$

$$Y \sim \text{Bernoulli}(\sigma(\mathbf{x}\beta))$$

where

$$\sigma(t) = \frac{1}{1 + e^{-t}} = \frac{e^t}{e^t + 1}$$

## Example 2: logistic regression



$$\sigma : \mathbb{R} \rightarrow (0, 1)$$

## Example 2: logistic regression

Claim:

$$X\beta = \log\left(\frac{p}{1-p}\right)$$

*The proof is left as an exercise.*

## Example 2: logistic regression

Some context for the data:

- ▶  $y$ : outcome of a soccer game.
- ▶  $x_1$ : goal difference between the home and the away team based on previous games *this season*.
- ▶  $x_2$ : goal difference between the home and the away team based on previous games *last season*.

Feel free to ask me more questions about the data.

## Example 2: logistic regression

*Exercise: write and fit a logistic regression.*

- ▶ Use `data/logistic.data.r`
- ▶ First use only one covariate, namely  $x_1$ .
- ▶ Use `y ~ bernoulli_logit(beta * x1)`.

## Example 2: logistic regression

- ▶ How should we do ppc's here?

## Example 2: logistic regression

- ▶ Can use misclassification rate.

```
generated quantities {  
  int y_pred[N];  
  int sum_err = 0;  
  
  for (i in 1:N)  
    y_pred[i] = bernoulli_logit_rng(beta * x1[i]);  
  for (i in 1:N) sum_err += (y[i] != y_pred[i]);  
}
```

- ▶ Can also compute the misclassification rate on a validation set.



## Example 2: logistic regression

```
$summary
      mean      se_mean      sd      2.5%      25%      50%
beta    0.5393828 0.003007621 0.1044100  0.3559608  0.4661776  0.5319302
lp__   -39.3343232 0.017393778 0.6806572 -41.2310970 -39.5035019 -39.0706867
sum_err 24.6530000 0.075131633 3.9385237 17.0000000 22.0000000 24.5000000
      75%      97.5%    n_eff    Rhat
beta    0.6055403  0.7611313 1205.140 1.005007
lp__   -38.8904949 -38.8440736 1531.331 1.000010
sum_err 27.0000000 33.0000000 2748.029 1.001644
```

- ▶ The 95th quantile interval for `sum_err` is [17, 33]
- ▶  $\beta$  is positive, which is what we would expect.

## Example 2: logistic regression

*Exercise: Augment the model by adding  $x_2$  as a covariate. Has the model improved?*

## Example 2: logistic regression

```
$summary
      mean      se_mean      sd      2.5%      25%      50%
beta[1] -0.4567048 0.008312798 0.3356420 -1.202829 -0.6677736 -0.4243604
beta[2]  1.6653154 0.015719100 0.5613639  0.729895  1.2450485  1.6352881
lp__    -5.5497786 0.026844508 0.9695979 -8.204500 -5.9291885 -5.2451473
sum_err  2.6777500 0.024126722 1.3432632  1.000000  2.0000000  2.0000000
      75%      97.5%      n_eff      Rhat
beta[1] -0.2165342 0.1250064 1630.265 1.0008829
beta[2]  2.0328909 2.8516718 1275.362 1.0027203
lp__    -4.8622523 -4.5788527 1304.585 1.0013099
sum_err  3.0000000 6.0000000 3099.742 0.9996283
```

- ▶ The 95th quantile interval for `sum_err` is [1, 6]
- ▶  $\beta_1$  is negative, which is surprising.
- ▶  $\beta_2$  is positive, which is what we expect.

## Example 2: logistic regression

Whether the model is “working” depends on our utility function.

- ▶ Do we care about predictions?
- ▶ Do we care about inference?

# Office hour

Use the remaining time for questions.

# References I

[Betancourt, 2018] Betancourt, M. (2018).

Towards a principled bayesian workflow.

[Gabry et al., 2017] Gabry, J., Simpson, D., Vehtari, A., Betancourt, M., and Gelman, A. (2017).

Visualization in bayesian workflow.

*Royal Journal of Statistics, section A*, 182:1 –14.